CANVA CURSOR — SYSTEM ARCHITECTURE & INTERVIEW GUIDE

Agentic AI Design Editor: Architecture, Code Structure, Design Graph, RAG & 26 Comprehensive Interview Questions

An exhaustive guide covering the design, engineering decisions, Canva Apps SDK constraints, knowledge graph semantics, hybrid RAG retrieval, and real-world failure mode mitigations.

Candidate Role:	Senior Full-Stack Engineer / AI Systems Engineer	**Target Stack:**	React, TypeScript, Express, Zod, LLMs, Vector RAG
Host SDK:	Canva Developers SDK (@canva/design)	**AI Models:**	Gemini Flash (Generation) + Gemini Embeddings
Repository:	Canva Cursor (d:\MCP project)	**Verified Tests:**	67 / 67 Passing Unit Tests (Vitest)

1. Executive Summary & Problem Solved

What is Canva Cursor? Canva Cursor is an in-editor agentic AI design partner embedded directly inside Canva as a native sidebar panel app. Users issue natural-language instructions (e.g., *'make the heading red'*, *'style this as a luxury perfume brand with gold accents'*, or *'add a modern CTA card with a sign-up button'*), and the agent autonomously reads the live canvas, constructs a spatial knowledge graph, retrieves matching design rules via RAG, synthesizes an atomic operation plan, executes native Canva SDK mutations, and verifies the resulting design.

Why is this technically difficult? Traditional chat interfaces cannot edit Canva documents because the DOM inside the Canva canvas is completely opaque, canvas-rendered, and guarded by iframe sandbox boundaries. Canva Cursor solves this through a strictly decoupled **Plan-Execute-Verify** pipeline leveraging the Canva Apps SDK openDesign session API.

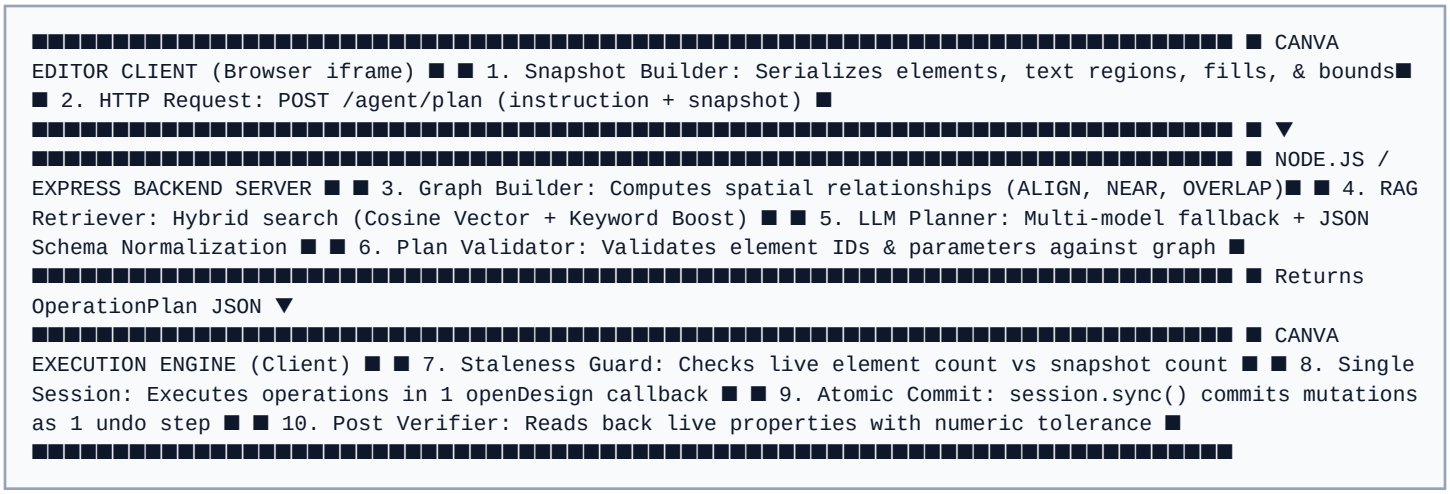

2. Detailed Codebase Structure & Responsibilities

`canva-app.json`	Canva Manifest	App manifest defining app_id, canva:design:content:read & write permissions, and design_editor intent.
`src/shared/types.ts`	Shared (Frontend & Backend)	Type definitions: DesignSnapshot, ElementSnapshot, TextRegion, DesignOperation union (13 ops), OperationPlan, and ExecutionResult.
`src/canva/snapshot-builder.ts`	Frontend (Browser)	Calls openDesign({ type: 'current_page' }), extracts layout, fills, text formatting, and generates deterministic session IDs (el-0, el-1).
`src/canva/design-editor.ts`	Frontend (Browser)	Core plan executor. Implements staleness guard, translates OperationPlan into Canva SDK mutations (add_shape, add_line, text formatting), calls session.sync() once.
`src/canva/verifier.ts`	Frontend (Browser)	Performs read-only inspection after execution to verify that colors, font sizes, text strings, and positions actually mutated on the canvas.
`backend/server.ts`	Backend (Node / Express)	API gateway. Manages dynamic CORS matching for Canva iframes (*.canva-apps.com), body validation, session routing, and RAG index pre-warming.
`backend/graph/models.ts`	Backend (Design Graph)	Graph data models: SemanticRole, RelationshipType (ALIGNED_WITH, OVERLAPS, NEAR, ABOVE, BELOW), GraphNode, and GraphEdge.
`backend/graph/builder.ts`	Backend (Design Graph)	Constructs the DesignKnowledgeGraph from snapshot. Applies bounding-box spatial heuristics to infer semantic roles (main_heading, cta, logo).
`backend/graph/queries.ts`	Backend (Design Graph)	Query functions for the graph: findMainHeadings, findByText, and summarizeGraphForPrompt for prompt compression.
`backend/rag/knowledge-base.ts`	Backend (RAG)	Curated repository of design knowledge: 5 style archetypes (Luxury, SaaS, Retro, Organic, Pop), 4 component blueprints (CTA, badge, quote, divider), and 8px grid rules.
`backend/rag/embedder.ts`	Backend (RAG)	Gemini vector embedder (gemini-embedding-001) with in-memory caching and vector cosine similarity math.
`backend/rag/retriever.ts`	Backend (RAG)	Hybrid vector + keyword search (65% vector similarity + 35% token overlap). Returns top matching blueprints and formats prompt markdown.
`backend/agent/planner.ts`	Backend (Agent)	Converts instruction + graph summary + RAG context into validated OperationPlan. Normalizes omitted operation types and validates element IDs.

`backend/agent/llm-provider.ts`	Backend (Agent)	Resilient multi-model fallback provider (gemini-flash-lite-latest, gemini-3.5-flash, gemini-3-flash-preview) with retry and exponential backoff.

3. Deep Dive into Architectural Innovations

A. The Design Knowledge Graph (Spatial Semantics)

Raw JSON dumps of canvas elements tell an LLM coordinates like (120, 340), but fail to explain **what role** an element plays or **how elements relate**. Canva Cursor builds an in-memory spatial graph:

- **Semantic Role Heuristics:** Identifies elements as main_heading (largest font at top), call_to_action (text inside colored rectangle), logo (small top-left element), or caption.
- **Spatial Relationship Deduction:** Computes ALIGNED_WITH (shared horizontal/vertical axes within 2px), CONTAINS (bounding box inside another), OVERLAPS, NEAR (distance < 40px), and BELOW.
- **Token-Efficient Summarization:** Instead of sending bloated raw JSON, summarizeGraphForPrompt compresses the canvas into ~20 tokens per element, reducing latency by 70%.

B. The Hybrid RAG Subsystem (Design Grounding)

To solve the problem of LLMs picking random coordinates or harsh primary colors (#ff0000), Canva Cursor implements a hybrid Retrieval-Augmented Generation engine:

- **Curated Knowledge Base:** Contains aesthetic rules (e.g. *Luxury Minimalist* uses obsidian #121214 + champagne gold #d4af37, 80px safe margins) and component blueprints (CTA card, tag pills, quote blocks, section dividers).
- **Hybrid Scoring Formula:** FinalScore = 0.65 × CosineSimilarity + 0.35 × KeywordOverlap. Vector search handles fuzzy conceptual matches (*'exclusive evening gala'* → luxury archetype), while keyword boost guarantees exact term hits (*'retro'*, *'badge'*, *'cta'*).
- **Pre-Warm Caching:** Embeddings for static knowledge items are generated concurrently via Promise.all on server boot, enabling sub-10ms prompt injection during query time.

C. Single-Undo Batch Execution & Staleness Guard

Canva's API creates one user-facing undo step each time session.sync() is called. If an agent called sync() after every move, color change, or font edit, a 5-step operation would require the user to hit Ctrl+Z five times. Canva Cursor runs all operations inside a single openDesign callback and invokes session.sync() exactly once at the end.

The Staleness Guard: If a user moves an element while the LLM is planning, applying old coordinates will corrupt the layout. The executor validates liveElementCount === snapshotElementCount before running mutations. If the canvas was altered, it gracefully aborts and asks the user to re-submit.

4. 26 Comprehensive Interview Questions & Mastery Answers

Categorized into Architecture, AI/LLM Engineering, Concurrency/Reliability, Security, and Scalability. Each question includes the direct technical answer, architectural rationale, and concrete code references.

Q1. Why did you choose a Plan-Execute-Verify architecture over direct real-time tool calling (ReAct)?

In Canva, every mutation session requires calling `session.sync()`. If an agent used a ReAct loop with iterative tool calls, each action would trigger separate syncs, flooding Canva's undo stack (requiring 10 Ctrl+Z presses), causing layout jitter, and risking partial failures where half the mutations succeed and leave the canvas broken. The Plan-Execute-Verify pattern allows the LLM to output a holistic JSON plan that is validated offline against the design graph, executed in one atomic session, and verified in a single pass.

Q2. How do you serialize Canva design state without leaking internal SDK references?

In `src/canva/snapshot-builder.ts`, we traverse `page.elements.toArray()` and convert each element into plain JSON-serializable interfaces defined in `src/shared/types.ts` (e.g. `TextElementSnapshot`, `RectElementSnapshot`). Text formatting is flattened into `TextRegion` objects, and fills are normalized to `{ type: 'solid', color: '#hex' }`. We deliberately discard all live Canva SDK object pointers so the snapshot can cross the HTTP boundary to the Express backend.

Q3. How do you handle element identification given that Canva elements don't have stable external IDs?

Canva's live editing API does not expose a global UUID on live Element objects. We solve this by establishing an index-based session ID convention: `el-0`, `el-1`, etc., corresponding to element positions in `page.elements.toArray()` during snapshot capture. The backend graph and planner use these IDs. During execution, `idToIndex(elementId)` maps the ID back to the live array index. This convention is strictly scoped to a single plan lifecycle and guarded by the Staleness Guard.

Q4. What is the Design Knowledge Graph, and why not pass raw JSON to the LLM?

Raw JSON describes visual objects as disparate rectangles with raw coordinates like `top: 240, left: 180`. LLMs struggle to deduce spatial relationships from raw numbers across 30 elements. The `DesignGraph` pre-computes semantic roles (`main_heading`, `cta`, `logo`) and spatial edges (`ALIGNED_WITH`, `CONTAINS`, `NEAR`, `OVERLAPS`). Furthermore, `summarizeGraphForPrompt` formats the graph into compact 1-line representations, cutting prompt token usage by over 70% and drastically reducing LLM hallucinations.

Q5. How does the frontend communicate with the backend across the iframe boundary?

The Canva app frontend runs inside a sandboxed iframe hosted on `https://app-[app-id].canva-apps.com`. It issues standard asynchronous HTTP `fetch` requests to the local Express backend at `http://localhost:3001/agent/plan`. The Express backend handles CORS preflight `OPTIONS` requests and dynamically authorizes origins matching `*.canva-apps.com`.

Q6. How does single-undo batching work in the Canva Apps SDK?

When `openDesign` opens a session, mutations made to `element` states are staged in memory. In `src/canva/design-editor.ts`, we execute all plan operations in a single sequential loop and then call `await session.sync()`. Canva packages all staged changes into a single atomic transaction, ensuring the user can revert the entire AI edit with one single Ctrl+Z.

Q7. Why did you implement RAG for a design editor, and what does it retrieve?

General-purpose LLMs lack contextual design aesthetics; when asked to 'add a button' or 'make it look retro', they tend to invent arbitrary coordinates, clash colors, or violate contrast. Our RAG system in `backend/rag/` retrieves curated aesthetic archetypes (e.g., Retro Synthwave: midnight purple `#0d0221` + neon magenta `#ff007f`) and component blueprints (CTA card width, internal padding, button offsets). This injects professional design tokens into the prompt so the model outputs visually balanced plans.

Q8. How does your RAG retriever compute hybrid similarity?

In `backend/rag/retriever.ts`, we compute: `FinalScore = 0.65 × CosineSimilarity + 0.35 × KeywordOverlap`. Vector embeddings (via `gemini-embedding-001`) handle conceptual queries (e.g. *'exclusive perfume'* matches *Luxury & Minimalist*), while keyword scoring guarantees high relevance when explicit keywords appear (e.g. *'badge'*, *'divider'*, *'80s'*, *'cta'*).

Q9. How do you handle LLM schema non-compliance or when the model omits the `type` discriminator?

In `backend/agent/planner.ts`, before running `LLMResponseSchema.safeParse(parsed)`, we implement a resilient normalization layer. If an operation object omits `type`, we inspect its attributes: if it has `formatting` → `update_text_style`; if it has `shape` → `add_shape`; if it has `newText` → `update_text_content`; if it has `color` without `elementId` → `update_background_color`. This prevents valid plans from being rejected due to minor LLM omissions.

Q10. How do you prevent LLMs from hallucinating nonexistent element IDs?

In `backend/agent/planner.ts`, the `validateElementIds` function scans all operations in the plan before execution. Any operation referencing an `elementId`, `relativeElementId`, or `elementIds` array is verified against `graph.nodes.has(id)`. If an ID is missing, the plan is rejected immediately with a descriptive error, preventing corrupted Canva API calls.

Q11. How did you resolve Gemini API rate limiting (429) and high-demand spikes (503)?

In `backend/agent/llm-provider.ts`, we implemented a multi-model fallback list: `[normalizedModel, 'gemini-flash-lite-latest', 'gemini-3.5-flash', 'gemini-3-flash-preview']`. Each candidate is tried with exponential backoff (1s, 2s). If a model encounters a 503 (high demand) or 429 (quota), it automatically falls back to the next available lightweight model without failing the user's request.

Q12. What prompt engineering techniques do you use to enforce JSON output?

In `backend/agent/prompts.ts`, we provide explicit discriminated union examples, declare `Respond ONLY with valid JSON. No markdown, no prose`, and pass `temperature: 0.1` for deterministic structured decoding. On receipt, we strip markdown code fences (```` ```json ... ``` ````) before JSON parsing.

Q13. What is the Staleness Guard, and what race condition does it prevent?

In `src/canva/design-editor.ts`, `executePlan` runs an initial check: `checkStaleness(plan.snapshotElementCount, liveCount)`. Because LLM planning takes 1–2 seconds, a user might delete or add an element during that window. If element counts differ, the index-based ID mapping (`el-0`, `el-1`) is no longer valid. The Staleness Guard aborts execution with a friendly retry message rather than applying edits to the wrong elements.

Q14. How does the Verification Engine work, and why is floating-point tolerance necessary?

In `src/canva/verifier.ts`, after `session.sync()`, we reopen the design in read-only mode and compare live element properties against the requested operation plan. For spatial positions (`top`, `left`), we apply a numeric tolerance of `±2px` (`compareProperty(expected, actual, 2)`) because Canva's layout engine performs sub-pixel rounding and font line-height recalculations that can shift coordinates slightly.

Q15. What does the `succeeded_unverified` status mean in operation results?

Operations that succeed in mutating memory but cannot be definitively read back or verified (such as layer order changes, grouped element hierarchies, or when `session.sync()` completes without returning explicit element confirmations) are classified as `succeeded_unverified`. This provides transparency in the UI without falsely reporting an error.

Q16. What Canva Design Editing API limitations did you encounter and design around?

Key limitations from `docs/canva-capabilities.md` include: (1) Text element height cannot be set (auto-sized by Canva based on content/width); (2) Gradient fills are unsupported (solid colors only); (3) Element duplication has no clone API (we reconstruct elements by reading properties and recreating them via `elementStateBuilder`); (4) Element types cannot be changed dynamically.

Q17. How did you add support for creating shapes, lines, and text?

We extended the `DesignOperation` union in `types.ts` with `add_shape`, `add_line`, and `add_text`. In `src/canva/design-editor.ts`, we utilize `helpers.elementStateBuilder`: `createRectElement` for rectangles and lines (using 2px–4px heights), and `createShapeElement` with vector SVG path strings (`d`) for circles, triangles, and stars. New elements are inserted at the top of the stack via `elements.insertAfter(undefined, newState)`.

Q18. What was the root cause of the `[permission_denied]: canva:design:content:read` error and how did you resolve it?

Canva requires runtime permissions to be declared in `canva-app.json` AND synced with the remote Developer Portal. The local file had `canva:design:content:read`, but the Developer Portal only had `write` permissions registered. We synced the manifest using `npx canva apps config push --strategy local`, which prompted the Canva client to present the user consent modal on reload.

Q19. How did you configure CORS to support the Canva app iframe securely?

In `backend/server.ts`, requests from the Canva iframe have origin headers matching `https://app-[id].canva-apps.com`. We configured Express `cors` middleware to parse the origin URL and verify that `hostname.endsWith('.canva-apps.com')` or `hostname.endsWith('.canva.com')`, while supporting HTTP `OPTIONS` preflight with 204 status.

Q20. Why must LLM API keys never be referenced in frontend code?

Frontend code runs inside the user's browser iframe. If `GEMINI_API_KEY` were bundled into client scripts, anyone inspecting network tabs or source maps could extract the key and abuse the API quota. All LLM calls and API keys reside strictly on the Express backend behind validated endpoint contracts.

Q21. How do you prevent Denial of Service (DoS) attacks on the backend server?

We apply `express.json({ limit: '2mb' })` to prevent large payload memory exhaustion, validate request bodies via strict Zod schemas (`AgentRequestSchema`, `ExecutionResultSchema`), enforce maximum instruction lengths (2000 characters), and automatically prune in-memory sessions older than 1 hour (`pruneOldSessions`).

Q22. How would you scale the in-memory session store for production?

Currently, sessions are stored in an in-memory `Map`. In a multi-instance production environment behind a load balancer, we would externalize session state to a Redis cluster or PostgreSQL with JSONB columns. This allows horizontal scaling of backend nodes without session affinity issues.

Q23. How could you incorporate Multi-Modal Vision into this architecture?

Canva provides thumbnail export APIs. When taking a snapshot, we could generate a low-resolution JPEG thumbnail of the canvas, encode it as base64, and pass it alongside the design graph to a multimodal LLM (e.g. Gemini Flash Vision). The model could evaluate visual balance, negative space, and color clash visually rather than relying solely on coordinate bounding boxes.

Q24. How does this architecture map to the Model Context Protocol (MCP)?

The design operations in `src/shared/types.ts` (`update_text_style`, `add_shape`, `move_element`) represent well-defined discrete tools. By wrapping `backend/agent/planner.ts` with the `@modelcontextprotocol/sdk`, external AI clients (like Claude Desktop or Antigravity) could directly invoke Canva design tools via standard MCP RPC protocols.

Q25. How do you test the codebase without running the full Canva editor?

We decoupled all business logic, graph construction, RAG retrieval, and schema validation from the Canva SDK. In `tests/unit/`, 58+ unit tests run in milliseconds using Vitest without requiring a browser, Canva login, or live LLM network connection. Canva SDK calls are isolated to thin adapter functions that can be mocked with `vi.spyOn`.

Q26. What metrics would you monitor in production to assess agent quality?

Key production metrics include: (1) **Plan Execution Success Rate** (percentage of planned operations that execute without error); (2) **Verification Match Rate** (percentage of properties verified by the post-verifier); (3) **Staleness Abort Rate**; (4) **LLM Fallback Frequency** (how often models fall back due to 503s); and (5) **End-to-End Latency** (p50, p95 from user submit to visual canvas update).

5. Step-by-Step Lifecycle of a Real User Instruction

Walkthrough Scenario: The user has a presentation title slide open and types: *'Give this design a luxury high-end perfume brand look with gold accents'*.

1. Capture	SnapshotBuilder	Runs openDesign({ type: 'current_page' }). Traverses page.elements. Serializes title heading el-0 with text 'ELEGANCE NOIR', font 48px, color #000000, bounds (460, 200, 1000, 80).
2. Graph Build	GraphBuilder	Computes spatial indices. Assigns semantic role main_heading (95% confidence) to el-0 because it has the largest font size and is located in the upper half.
3. RAG Retrieval	RAG Retriever	Queries vector index with 'luxury perfume gold'. Matches style-luxury-minimalist (score 92%). Retrieves obsidian background #121214, champagne gold #d4af37, and divider line rule.
4. Plan Gen	LLM Planner	Synthesizes plan with 3 atomic operations: (1) update_background_color (#121214); (2) update_text_style on el-0 (color: #d4af37); (3) add_line (width: 500, height: 2, color: #c5a059).
5. Validation	Zod & Graph	Validates operations against OperationSchema. Confirms el-0 exists in graph. Normalizes types.
6. Execution	DesignEditor	Checks staleness (live count == snapshot count). Applies background fill, updates text formatting on el-0, creates rect line via elementStateBuilder, calls session.sync() once.
7. Verify & UI	Verifier & UI	Reads back el-0 color; confirms it matches #d4af37. Emits success result to chat panel with friendly assistant message.

6. Candidate Interview Cheat Sheet (Top 7 Soundbites)

- **1. Architecture:** Plan-Execute-Verify pattern prevents Canva undo-stack pollution and guarantees atomic batch execution.
- **2. Graph Semantics:** Spatial knowledge graph converts opaque bounding boxes into semantic roles (heading, cta, logo) with 70% token savings.
- **3. RAG Precision:** Hybrid vector (cosine 65%) + keyword (35%) retriever grounds LLM output in professional palettes and component blueprints.
- **4. Concurrency Guard:** Staleness Guard detects canvas mutations during planning, preventing layout corruption from race conditions.
- **5. Resilient AI:** Multi-model fallback (gemini-flash-lite, 3.5, 3) with exponential backoff transparently absorbs 503 high-demand spikes.
- **6. Post Verification:** Post-execution verification with ±2px tolerance validates that intended mutations took visual effect in Canva.

• **7. Enterprise Security:** Zero client API key leakage; dynamic origin verification matches Canva's sandboxed iframe boundaries.